

# Internship Task - 2

## Install and import the Lib

```
In [1]: # !pip install -q kaggle
# !pip install -U keras-efficientnet-v2 -q
```

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import os
import tensorflow as tf
import tensorflow_hub as hub
import cv2
from sklearn.model_selection import train_test_split
# import keras_efficientnet_v2
from tensorflow.keras.applications import efficientnet
from tensorflow.keras.models import Sequential, load_model
from tensorflow.keras.layers import Dense, Activation, Conv2D, Flatten, Dropout
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from keras import regularizers
from tensorflow.keras import metrics, optimizers, Sequential, activations, constraints
from sklearn.metrics import f1_score, classification_report, confusion_matrix
import keras.backend as K
from tensorflow import keras
import warnings
warnings.filterwarnings('ignore')
```

## Getting the data

```
In [3]: # !rm -r ~/.kaggle
# !mkdir ~/.kaggle
# !mv ./kaggle.json ~/.kaggle/
# !chmod 600 ~/.kaggle/kaggle.json
```

```
In [4]: # !kaggle datasets download lexset/synthetic-asl-alphabet
```

```
In [5]: # !unzip -qq "synthetic-asl-alphabet.zip" -d "synthetic-asl-alphabet"
```

```
In [6]: train_dir = '/kaggle/input/synthetic-asl-alphabet/Train_Alphabet'
test_dir = '/kaggle/input/synthetic-asl-alphabet/Test_Alphabet'
```

```
In [7]: def load_df(dir_path):
    df = {
        'filename': [],
        'label': []
    }
```

```

classes = os.listdir(dir_path)
for c in classes:
    path = dir_path + '/' + c

    img_paths = os.listdir(path)

    for img_path in img_paths:
        df['filename'].append(c + '/' + img_path)
        df['label'].append(c)

df = pd.DataFrame(df)
return df

train_df = load_df(train_dir)
test_df = load_df(test_dir)

```

In [8]: `train_df.sample(5)`

Out[8]:

|              | filename  | label |
|--------------|---|-------|
| <b>8944</b>  | X/8e5d20e5-109c-4bd7-9e66-8f0d1b16080f.rgb_000... | X     |
| <b>6525</b>  | U/998bfaa6-0fd4-495a-b66c-0e522864baca.rgb_000... | U     |
| <b>4929</b>  | H/885c0cf6-84e0-4cfa-ad2e-523a4d7b4467.rgb_000... | H     |
| <b>8918</b>  | X/a94d77e6-1fdf-4fc0-bc7a-0a04e15db8d5.rgb_000... | X     |
| <b>10763</b> | Q/0d99d487-da2f-4fc8-9521-a88dcd75283c.rgb_000... | Q     |

## Split train-valid

In [9]:

```

X = train_df['filename']
y = train_df['label']

X_train, X_valid, y_train, y_valid = train_test_split(X, y, stratify=y, random_state=42)
print(X_train.shape)
print(X_valid.shape)

```

```

(19440,)
(4860,)

```

In [10]:

```

train = pd.concat([X_train, y_train], axis=1)
valid = pd.concat([X_valid, y_valid], axis=1)
test = test_df

```

In [11]: `train['label'].value_counts()`

```
Out[11]: E      720
          K      720
          O      720
          M      720
          X      720
          P      720
          N      720
          B      720
          W      720
          S      720
          Z      720
          V      720
          G      720
          D      720
          Y      720
          J      720
          L      720
          U      720
          A      720
          C      720
          R      720
          H      720
          Blank  720
          I      720
          T      720
          F      720
          Q      720
          Name: label, dtype: int64
```

```
In [12]: valid['label'].value_counts()
```

```
Out[12]: C      180
         B      180
         A      180
         G      180
         X      180
         I      180
         K      180
         Y      180
         P      180
         H      180
         T      180
         O      180
         U      180
         E      180
         J      180
         L      180
         Q      180
         V      180
         M      180
         R      180
         S      180
         F      180
         D      180
         W      180
         Blank  180
         N      180
         Z      180
         Name: label, dtype: int64
```

```
In [13]: test['label'].value_counts()
```

```
Out[13]: N      100
         S      100
         D      100
         W      100
         L      100
         P      100
         C      100
         Z      100
         V      100
         T      100
         O      100
         A      100
         G      100
         Blank  100
         R      100
         Y      100
         Q      100
         K      100
         X      100
         M      100
         U      100
         E      100
         H      100
         F      100
         I      100
         B      100
         J      100
         Name: label, dtype: int64
```

```
In [14]: IMAGE_SIZE = (224, 224)

datagen = ImageDataGenerator(
    rescale=1./255,
    # rotation_range=40,
    zoom_range=0.2,
    width_shift_range=0.05,
    height_shift_range=0.05,
    # brightness_range=[0.5, 1.0],
)

train_generator = datagen.flow_from_dataframe(
    dataframe = train,
    directory = train_dir,
    x_col='filename',
    y_col='label',
    batch_size=32,
    # color_mode="rgb",
    seed=42,
    shuffle=True,
    class_mode='sparse',
    target_size=IMAGE_SIZE
)

valid_generator=datagen.flow_from_dataframe(
    dataframe=valid,
    directory=train_dir,
```

```

        x_col="filename",
        y_col="label",
        batch_size=32,
        # color_mode="rgb",
        seed=42,
        shuffle=True,
        class_mode="sparse",
        target_size=IMAGE_SIZE
    )

test_datagen=ImageDataGenerator(rescale=1./255.)

test_generator=test_datagen.flow_from_dataframe(
    dataframe=test,
    directory=test_dir,
    x_col="filename",
    y_col=None,
    batch_size=1,
    # color_mode="rgb",
    seed=42,
    shuffle=False,
    class_mode=None,
    target_size=IMAGE_SIZE
)

```

Found 19440 validated image filenames belonging to 27 classes.  
 Found 4860 validated image filenames belonging to 27 classes.  
 Found 2700 validated image filenames.

## EfficientNetV2B0

### Training

```

In [15]: tf.keras.backend.clear_session()
num_classes = 27

backbone = tf.keras.applications.EfficientNetB0(include_top=False)
backbone.trainable = False

inputs = Input(shape=IMAGE_SIZE + (3,), name='inputs')
x = backbone(inputs, training=False)
x = GlobalAveragePooling2D()(x)
x = Dropout(0.2)(x)
# x = Dense(100, activation='relu')(x)
outputs = Dense(27, activation='softmax')(x)

model_efficientnet = tf.keras.Model(inputs=inputs, outputs=outputs, name='ef
model_efficientnet.summary()

```

```
2023-01-06 08:32:48.777776: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2023-01-06 08:32:48.889253: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2023-01-06 08:32:48.890088: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2023-01-06 08:32:48.891464: I tensorflow/core/platform/cpu_feature_guard.cc:142] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: AVX2 AVX512F FMA
To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.
2023-01-06 08:32:48.891823: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2023-01-06 08:32:48.892580: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2023-01-06 08:32:48.893229: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2023-01-06 08:32:51.147747: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2023-01-06 08:32:51.148606: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2023-01-06 08:32:51.149356: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2023-01-06 08:32:51.149953: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1510] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 15401 MB memory: -> device: 0, name: Tesla P100-PCIE-16GB, pci bus id: 0000:00:04.0, compute capability: 6.0
```

Downloading data from [https://storage.googleapis.com/keras-applications/efficientnetb0\\_notop.h5](https://storage.googleapis.com/keras-applications/efficientnetb0_notop.h5)  
 16711680/16705208 [=====] - 0s 0us/step  
 16719872/16705208 [=====] - 0s 0us/step  
 Model: "efficientnet\_model"

| Layer (type)                    | Output Shape             | Param # |
|---------------------------------|--------------------------|---------|
| inputs (InputLayer)             | [(None, 224, 224, 3)]    | 0       |
| efficientnetb0 (Functional)     | (None, None, None, 1280) | 4049571 |
| global_average_pooling2d (Gl    | (None, 1280)             | 0       |
| dropout (Dropout)               | (None, 1280)             | 0       |
| dense (Dense)                   | (None, 27)               | 34587   |
| =====                           |                          |         |
| Total params: 4,084,158         |                          |         |
| Trainable params: 34,587        |                          |         |
| Non-trainable params: 4,049,571 |                          |         |

```

In [16]: tf.keras.backend.clear_session()
num_classes = 27
model_url = 'https://tfhub.dev/google/imagenet/efficientnet_v2_imagenet1k_b0

feature_extractor_layer = hub.KerasLayer(model_url,
                                          trainable=False, # freeze the under
                                          name='feature_extraction_layer',
                                          input_shape=IMAGE_SIZE + (3,))

# Create our own model
# inputs = Input(shape=IMAGE_SIZE + (3,), name='inputs')
# x = feature_extractor_layer(inputs, training=False)
# x = Dropout(0.2)(x)
# outputs = Dense(27, activation='softmax')(x)

# model_efficientnet = tf.keras.Model(inputs=inputs, outputs=outputs, name=

model_efficientnet = tf.keras.Sequential([
    feature_extractor_layer, # use the feature extraction layer as the base
    Dropout(0.2),
    Dense(num_classes, activation='softmax') # create our own output layer
])

model_efficientnet.summary()
  
```



Model: "sequential"

| Layer (type)                    | Output Shape | Param # |
|---------------------------------|--------------|---------|
| =====                           |              |         |
| feature_extraction_layer (Ke    | (None, 1280) | 5919312 |
| =====                           |              |         |
| dropout (Dropout)               | (None, 1280) | 0       |
| =====                           |              |         |
| dense (Dense)                   | (None, 27)   | 34587   |
| =====                           |              |         |
| Total params: 5,953,899         |              |         |
| Trainable params: 34,587        |              |         |
| Non-trainable params: 5,919,312 |              |         |

```
In [17]: model_efficientnet.compile(optimizer=keras.optimizers.Adam(),
                                     loss='sparse_categorical_crossentropy',
                                     metrics=["accuracy"])
```

```
In [18]: cb = [
#         callbacks.EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True),
        callbacks.ModelCheckpoint('model_efficientnet.h5', save_best_only=True,
    ]

history_efficientnet = model_efficientnet.fit(
    train_generator,
    validation_data = valid_generator,
    epochs = 20,
    batch_size = 32,
    verbose = 1,
    callbacks = cb
)
```

```
2023-01-06 08:33:01.502070: I tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:185] None of the MLIR Optimization Passes are enabled (registered 2)
```

Epoch 1/20

```
2023-01-06 08:33:09.749165: I tensorflow/stream_executor/cuda/cuda_dnn.cc:369] Loaded cuDNN version 8005
```

608/608 [=====] - 714s 1s/step - loss: 1.7488 - accuracy: 0.5698 - val\_loss: 1.0928 - val\_accuracy: 0.7626  
Epoch 2/20  
608/608 [=====] - 493s 811ms/step - loss: 0.9571 - accuracy: 0.7684 - val\_loss: 0.8192 - val\_accuracy: 0.8109  
Epoch 3/20  
608/608 [=====] - 491s 808ms/step - loss: 0.7490 - accuracy: 0.8134 - val\_loss: 0.6701 - val\_accuracy: 0.8451  
Epoch 4/20  
608/608 [=====] - 499s 821ms/step - loss: 0.6413 - accuracy: 0.8355 - val\_loss: 0.6055 - val\_accuracy: 0.8486  
Epoch 5/20  
608/608 [=====] - 487s 801ms/step - loss: 0.5730 - accuracy: 0.8519 - val\_loss: 0.5650 - val\_accuracy: 0.8562  
Epoch 6/20  
608/608 [=====] - 486s 799ms/step - loss: 0.5229 - accuracy: 0.8580 - val\_loss: 0.5226 - val\_accuracy: 0.8599  
Epoch 7/20  
608/608 [=====] - 481s 791ms/step - loss: 0.4897 - accuracy: 0.8673 - val\_loss: 0.4817 - val\_accuracy: 0.8702  
Epoch 8/20  
608/608 [=====] - 481s 791ms/step - loss: 0.4586 - accuracy: 0.8751 - val\_loss: 0.4687 - val\_accuracy: 0.8778  
Epoch 9/20  
608/608 [=====] - 477s 785ms/step - loss: 0.4399 - accuracy: 0.8790 - val\_loss: 0.4552 - val\_accuracy: 0.8735  
Epoch 10/20  
608/608 [=====] - 483s 795ms/step - loss: 0.4202 - accuracy: 0.8825 - val\_loss: 0.4390 - val\_accuracy: 0.8819  
Epoch 11/20  
608/608 [=====] - 479s 788ms/step - loss: 0.4080 - accuracy: 0.8859 - val\_loss: 0.4270 - val\_accuracy: 0.8842  
Epoch 12/20  
608/608 [=====] - 493s 810ms/step - loss: 0.3968 - accuracy: 0.8836 - val\_loss: 0.4192 - val\_accuracy: 0.8815  
Epoch 13/20  
608/608 [=====] - 486s 799ms/step - loss: 0.3824 - accuracy: 0.8870 - val\_loss: 0.3977 - val\_accuracy: 0.8835  
Epoch 14/20  
608/608 [=====] - 502s 826ms/step - loss: 0.3651 - accuracy: 0.8942 - val\_loss: 0.3997 - val\_accuracy: 0.8866  
Epoch 15/20  
608/608 [=====] - 485s 798ms/step - loss: 0.3547 - accuracy: 0.8973 - val\_loss: 0.3747 - val\_accuracy: 0.8930  
Epoch 16/20  
608/608 [=====] - 489s 804ms/step - loss: 0.3511 - accuracy: 0.8989 - val\_loss: 0.3943 - val\_accuracy: 0.8877  
Epoch 17/20  
608/608 [=====] - 488s 803ms/step - loss: 0.3439 - accuracy: 0.8986 - val\_loss: 0.3714 - val\_accuracy: 0.8971  
Epoch 18/20  
608/608 [=====] - 497s 818ms/step - loss: 0.3367 - accuracy: 0.9012 - val\_loss: 0.3806 - val\_accuracy: 0.8934  
Epoch 19/20  
608/608 [=====] - 499s 820ms/step - loss: 0.3271 - accuracy: 0.9013 - val\_loss: 0.3700 - val\_accuracy: 0.8959


Epoch 20/20  
608/608 [=====] - 520s 856ms/step - loss: 0.3302 - accuracy: 0.9019 - val\_loss: 0.3700 - val\_accuracy: 0.8942

```
In [19]: fig, ax = plt.subplots(nrows = 1, ncols = 2, figsize = (12,6))

ax[0].plot(history_efficientnet.history['accuracy'], '-', label = 'Training accuracy')
ax[0].plot(history_efficientnet.history['val_accuracy'], '-', label = 'Validation accuracy')
ax[0].set_xlabel('Epochs')
ax[0].set_ylabel('Accuracy')
ax[0].set_title('Epochs & Training Accuracy', fontsize=20)
ax[0].legend(loc='best')

ax[1].plot(history_efficientnet.history['loss'], '-', label = 'Training loss')
ax[1].plot(history_efficientnet.history['val_loss'], '-', label = 'Validation loss')
ax[1].set_xlabel('Epochs')
ax[1].set_ylabel('loss')
ax[1].set_title('Epochs & Training loss', fontsize=20)
ax[1].legend(loc='best')

plt.show()
```



```
In [20]: model_loaded = model_efficientnet
```

```
In [21]: test_generator.reset()
pred = model_loaded.predict(test_generator, verbose=1)
```

2700/2700 [=====] - 65s 24ms/step

```
In [22]: pred = np.argmax(pred,axis=1)
labels = (train_generator.class_indices)
labels = dict((v,k) for k,v in labels.items())
pred = [labels[k] for k in pred]

print(accuracy_score(test['label'], pred))
print(classification_report(test['label'], pred))
```

0.9222222222222223

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| A            | 0.94      | 0.97   | 0.96     | 100     |
| B            | 0.99      | 0.92   | 0.95     | 100     |
| Blank        | 1.00      | 0.98   | 0.99     | 100     |
| C            | 0.99      | 0.98   | 0.98     | 100     |
| D            | 0.94      | 0.88   | 0.91     | 100     |
| E            | 0.91      | 0.94   | 0.93     | 100     |
| F            | 0.94      | 0.94   | 0.94     | 100     |
| G            | 0.91      | 0.96   | 0.93     | 100     |
| H            | 0.98      | 1.00   | 0.99     | 100     |
| I            | 0.91      | 0.89   | 0.90     | 100     |
| J            | 0.86      | 0.93   | 0.89     | 100     |
| K            | 0.94      | 0.88   | 0.91     | 100     |
| L            | 0.98      | 0.94   | 0.96     | 100     |
| M            | 0.80      | 0.90   | 0.85     | 100     |
| N            | 0.87      | 0.82   | 0.85     | 100     |
| O            | 0.97      | 0.99   | 0.98     | 100     |
| P            | 0.92      | 0.95   | 0.94     | 100     |
| Q            | 0.94      | 0.82   | 0.88     | 100     |
| R            | 0.96      | 0.94   | 0.95     | 100     |
| S            | 0.99      | 0.80   | 0.88     | 100     |
| T            | 0.85      | 0.97   | 0.91     | 100     |
| U            | 0.88      | 0.94   | 0.91     | 100     |
| V            | 0.81      | 0.90   | 0.85     | 100     |
| W            | 0.92      | 0.87   | 0.89     | 100     |
| X            | 0.84      | 0.92   | 0.88     | 100     |
| Y            | 0.97      | 0.96   | 0.96     | 100     |
| Z            | 0.97      | 0.91   | 0.94     | 100     |
| accuracy     |           |        | 0.92     | 2700    |
| macro avg    | 0.93      | 0.92   | 0.92     | 2700    |
| weighted avg | 0.93      | 0.92   | 0.92     | 2700    |

## Fine-tuning

```
In [23]: feature_extractor_layer.trainable = True

model_efficientnet.compile(optimizer=keras.optimizers.Adam(learning_rate=1e-4,
    loss='sparse_categorical_crossentropy',
    metrics=["accuracy"]))
```

```
In [24]: cb = [
    #     callbacks.EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True),
    callbacks.ModelCheckpoint('model_efficientnet_finetuned.h5', save_best_only=True),
]

train_generator.reset()
valid_generator.reset()

history_efficientnet_2 = model_efficientnet.fit(
    train_generator,
```

```

validation_data = valid_generator,
epochs = 10,
batch_size = 32,
verbose = 1,
callbacks = cb
)

```

```

Epoch 1/10
608/608 [=====] - 518s 834ms/step - loss: 1.2017 -
accuracy: 0.6690 - val_loss: 0.2236 - val_accuracy: 0.9292
Epoch 2/10
608/608 [=====] - 526s 866ms/step - loss: 0.3864 -
accuracy: 0.8756 - val_loss: 0.0874 - val_accuracy: 0.9753
Epoch 3/10
608/608 [=====] - 554s 911ms/step - loss: 0.2046 -
accuracy: 0.9336 - val_loss: 0.0442 - val_accuracy: 0.9858
Epoch 4/10
608/608 [=====] - 534s 879ms/step - loss: 0.1134 -
accuracy: 0.9632 - val_loss: 0.0265 - val_accuracy: 0.9922
Epoch 5/10
608/608 [=====] - 527s 867ms/step - loss: 0.0789 -
accuracy: 0.9738 - val_loss: 0.0195 - val_accuracy: 0.9947
Epoch 6/10
608/608 [=====] - 525s 864ms/step - loss: 0.0608 -
accuracy: 0.9801 - val_loss: 0.0158 - val_accuracy: 0.9953
Epoch 7/10
608/608 [=====] - 526s 864ms/step - loss: 0.0431 -
accuracy: 0.9857 - val_loss: 0.0119 - val_accuracy: 0.9971
Epoch 8/10
608/608 [=====] - 516s 848ms/step - loss: 0.0322 -
accuracy: 0.9898 - val_loss: 0.0069 - val_accuracy: 0.9984
Epoch 9/10
608/608 [=====] - 524s 860ms/step - loss: 0.0283 -
accuracy: 0.9917 - val_loss: 0.0072 - val_accuracy: 0.9988
Epoch 10/10
608/608 [=====] - 509s 836ms/step - loss: 0.0207 -
accuracy: 0.9935 - val_loss: 0.0067 - val_accuracy: 0.9984

```

```

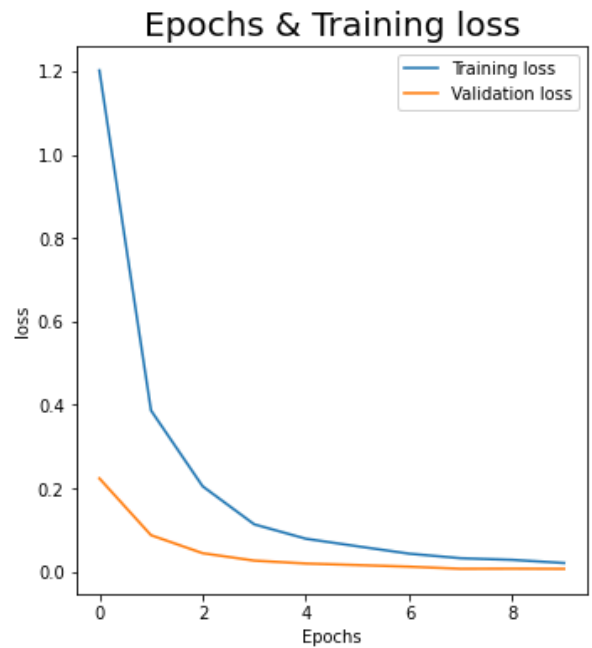
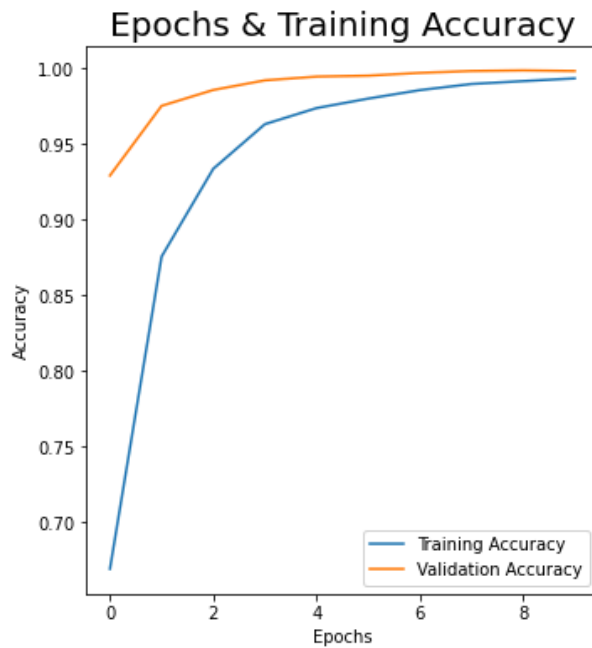
In [25]: fig, ax = plt.subplots(nrows = 1, ncols = 2, figsize = (12,6))

ax[0].plot(history_efficientnet_2.history['accuracy'], '-', label = 'Training accuracy')
ax[0].plot(history_efficientnet_2.history['val_accuracy'], '-', label = 'Validation accuracy')
ax[0].set_xlabel('Epochs')
ax[0].set_ylabel('Accuracy')
ax[0].set_title('Epochs & Training Accuracy', fontsize=20)
ax[0].legend(loc='best')

ax[1].plot(history_efficientnet_2.history['loss'], '-', label = 'Training loss')
ax[1].plot(history_efficientnet_2.history['val_loss'], '-', label = 'Validation loss')
ax[1].set_xlabel('Epochs')
ax[1].set_ylabel('loss')
ax[1].set_title('Epochs & Training loss', fontsize=20)
ax[1].legend(loc='best')

plt.show()

```



```
In [26]: model_loaded = model_efficientnet
```

```
In [27]: test_generator.reset()
pred = model_loaded.predict(test_generator, verbose=1)
```

2700/2700 [=====] - 51s 19ms/step

```
In [28]: pred = np.argmax(pred,axis=1)
labels = (train_generator.class_indices)
labels = dict((v,k) for k,v in labels.items())
pred = [labels[k] for k in pred]

print(accuracy_score(test['label'], pred))
print(classification_report(test['label'], pred))
```

0.9992592592592593

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| A            | 1.00      | 1.00   | 1.00     | 100     |
| B            | 1.00      | 1.00   | 1.00     | 100     |
| Blank        | 1.00      | 1.00   | 1.00     | 100     |
| C            | 1.00      | 1.00   | 1.00     | 100     |
| D            | 1.00      | 1.00   | 1.00     | 100     |
| E            | 1.00      | 1.00   | 1.00     | 100     |
| F            | 1.00      | 1.00   | 1.00     | 100     |
| G            | 0.98      | 1.00   | 0.99     | 100     |
| H            | 1.00      | 1.00   | 1.00     | 100     |
| I            | 1.00      | 1.00   | 1.00     | 100     |
| J            | 1.00      | 0.99   | 0.99     | 100     |
| K            | 1.00      | 1.00   | 1.00     | 100     |
| L            | 1.00      | 1.00   | 1.00     | 100     |
| M            | 1.00      | 1.00   | 1.00     | 100     |
| N            | 1.00      | 1.00   | 1.00     | 100     |
| O            | 1.00      | 1.00   | 1.00     | 100     |
| P            | 1.00      | 1.00   | 1.00     | 100     |
| Q            | 1.00      | 0.99   | 0.99     | 100     |
| R            | 1.00      | 1.00   | 1.00     | 100     |
| S            | 1.00      | 1.00   | 1.00     | 100     |
| T            | 1.00      | 1.00   | 1.00     | 100     |
| U            | 1.00      | 1.00   | 1.00     | 100     |
| V            | 1.00      | 1.00   | 1.00     | 100     |
| W            | 1.00      | 1.00   | 1.00     | 100     |
| X            | 1.00      | 1.00   | 1.00     | 100     |
| Y            | 1.00      | 1.00   | 1.00     | 100     |
| Z            | 1.00      | 1.00   | 1.00     | 100     |
| accuracy     |           |        | 1.00     | 2700    |
| macro avg    | 1.00      | 1.00   | 1.00     | 2700    |
| weighted avg | 1.00      | 1.00   | 1.00     | 2700    |